

Assignment 1: Word Representation, Text Classification, Machine Translation, and Pre-Trained Transformers

Part A:

1. Downloading the Corpus

Sanity Check

Note that the corpus size is 16498 when run in Jupyter and 16463 in Colab.

This training corpus contains 16498 sentences. The following print statement should return 16498.

```
[2] print(len(austen))
```

```
16463
```

```
[3] austen[0]
```

```
['[', 'Sense', 'and', 'Sensibility', 'by', 'Jane', 'Austen', '1811', '']']
```

2. Preprocessing the Training Corpus

Sanity Check

```
[['[', 'Sense', 'and', 'Sensibility', 'by', 'Jane', 'Austen', '1811', '']]
The new length of the preprocessed output: 12498
```

```
normalized_corpus[0]
```

```
['sense', 'sensibility', 'jane', 'austen']
```

```
sample = austen[:2] + austen[100:102]
preprocessed_sample = preprocess_corpus(sample)
```

```
print(len(sample), sample)
print()
print(len(preprocessed_sample), preprocessed_sample)
```

```
[['[', 'Sense', 'and', 'Sensibility', 'by', 'Jane', 'Austen', '1811', '']]
4 [['[', 'Sense', 'and', 'Sensibility', 'by', 'Jane', 'Austen', '1811', '']], ['CHAPTER', '1'], ['But', ' ', 'then', ' ', 'if', 'Mrs', ' ', 'Dashwood', 'should', 'live', 'fifteen', 'y
2 [['sense', 'sensibility', 'jane', 'austen'], ['mrs', 'dashwood', 'live', 'fifteen', 'years', 'shall', 'completely', 'taken']]
```

3. Creating the Corpus Vocabulary and Preparing the Data

Sanity Check

```
[9] print("Number of unique words:", len(word2idx))
```

```
Number of unique words: 10040
```

```
print("\nSample word2idx:", list(word2idx.items())[:10])
```

```
Sample word2idx: [(('engaging', 1), ('coldly', 2), ('excellencies', 3), ('sweetest', 4), ('cassino', 5), ('fulfil', 6), ('accumulation', 7), ('solidly', 8), ('natural', 9), ('redeem', 10))]
```

```
[11] print("\nSample idx2word:", list(idx2word.items())[:10])
```

```
Sample idx2word: [(1, 'engaging'), (2, 'coldly'), (3, 'excellencies'), (4, 'sweetest'), (5, 'cassino'), (6, 'fulfil'), (7, 'accumulation'), (8, 'solidly'), (9, 'natural'), (10, 'redeem')]
```

```
[12] print("\nSample sents_as_id:", prepareSentsAsId(preprocessed_sample))
```

```
Sample sents_as_id: [[438, 6885, 5070, 6141], [2053, 7434, 6231, 5231, 2686, 2628, 7499, 8743]]
```

4. Generating training instances

Sanity Check

```

(jane (5070), everybody (5172)) -> 0
(jane (5070), sense (438)) -> 1
(sensibility (6885), sense (438)) -> 1
(sensibility (6885), doors (1098)) -> 0
(sensibility (6885), rapturous (9994)) -> 0
(sense (438), jane (5070)) -> 1
(sensibility (6885), austen (6141)) -> 1
(sense (438), unanswerable (9839)) -> 0
(austen (6141), sense (438)) -> 1
(jane (5070), palanquins (9773)) -> 0
(austen (6141), jane (5070)) -> 1
(austen (6141), sensibility (6885)) -> 1
(jane (5070), sensibility (6885)) -> 1
(austen (6141), tumbling (1797)) -> 0
(jane (5070), stationing (2058)) -> 0
(austen (6141), conjurer (3176)) -> 0
(sense (438), ing (216)) -> 0
(sense (438), austen (6141)) -> 1
(sensibility (6885), jane (5070)) -> 1
(sense (438), lucy (9373)) -> 0
(austen (6141), amongst (7346)) -> 0
(sense (438), sensibility (6885)) -> 1
(jane (5070), austen (6141)) -> 1
(sensibility (6885), share (7124)) -> 0

```

5. Building the Skip-gram Neural Network Architecture

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 1)]	0	[]
input_4 (InputLayer)	[(None, 1)]	0	[]
target_embed_layer (Embedding)	(None, 1, 100)	1004100	['input_3[0][0]']
context_embed_layer (Embedding)	(None, 1, 100)	1004100	['input_4[0][0]']
reshape_2 (Reshape)	(None, 100)	0	['target_embed_layer[0][0]']
reshape_3 (Reshape)	(None, 100)	0	['context_embed_layer[0][0]']
dot_1 (Dot)	(None, 1)	0	['reshape_2[0][0]', 'reshape_3[0][0]']
dense_1 (Dense)	(None, 1)	2	['dot_1[0][0]']
=====			
Total params: 2,008,202			
Trainable params: 2,008,202			
Non-trainable params: 0			

6. Training the Model

Q1: What would the inputs and outputs to the model be?

Input is a one-hot vector representing the word we are inputting.

Output of a window size of k is $2k$ values distributions (i.e., $2k$ output neurons). Then, for each position, we get a probability distribution for each word in the vocabulary. Then, we select the word which has the highest probability.

Q2: How would you use the Keras framework to create this architecture?

Assuming a one-hot vector input, we have an Input layer

Then use an Embedding layer

Flatten the output of the Embedding layer

Then use a Dense layer with softmax activation

Q3: What are the reasons this training approach is considered inefficient?

There are a lot weights that need to be updated and as such a lot of data is required. Also the weights of non activated neurons are updated on each iteration, which is computationally inefficient. So, a better way is to use negative sampling and only updating a subset of relevant weights in the training.

7. Getting the Word Embeddings

	0	1	2	3	4	5	\
engaging	0.017765	0.000005	0.022112	-0.020767	-0.013893	-0.016174	
coldly	-0.027600	0.020438	0.049420	0.010220	0.023247	0.017857	
excellencies	-0.018321	-0.004841	0.032120	0.032563	0.006279	0.017992	
sweetest	-0.008631	0.026240	0.020887	0.020410	-0.013650	0.018767	
cassino	-0.012359	0.018434	0.027713	0.039947	0.016284	0.025528	
fulfil	-0.002769	0.006105	0.027198	-0.008085	-0.007095	-0.008833	
accumulation	-0.030529	0.014008	0.019657	0.021920	0.016112	0.012265	
solidly	-0.031012	0.019528	-0.008223	0.018680	0.023148	-0.009726	
natural	-0.014126	0.018897	-0.003177	0.005826	0.029550	0.023559	
redeem	-0.140309	-0.016170	0.112158	-0.064119	0.037389	-0.008252	

	6	7	8	9	...	90	91	\
engaging	0.012935	-0.023116	-0.002303	0.002457	...	-0.021339	0.005936	
coldly	-0.023826	0.012182	0.016692	-0.004947	...	-0.020355	-0.023555	
excellencies	-0.020537	0.006367	0.008557	-0.030948	...	-0.010404	-0.036386	
sweetest	-0.014266	0.017115	0.010169	0.002724	...	0.001609	0.003035	
cassino	-0.019760	0.017783	0.011593	-0.051846	...	-0.039112	0.000120	
fulfil	-0.011802	0.012419	0.014846	-0.026387	...	-0.037336	-0.023763	
accumulation	-0.019975	0.022388	-0.005133	-0.016820	...	0.006897	0.004993	
solidly	-0.030244	0.027006	-0.014429	-0.017336	...	0.002122	0.009850	
natural	-0.027112	0.008908	0.012467	-0.029214	...	-0.035930	-0.000452	
redeem	-0.022950	-0.047449	0.045108	-0.050998	...	-0.032899	0.064864	

	92	93	94	95	96	97	\
engaging	-0.008643	0.005085	0.014789	-0.008110	0.014664	0.008601	
coldly	0.064867	0.025325	-0.036700	-0.035499	0.051797	0.026862	
excellencies	-0.007756	0.018761	-0.027996	0.002799	0.029025	0.003498	
sweetest	0.033030	0.020545	0.004448	-0.007108	0.013759	0.013011	
cassino	0.043426	0.044663	-0.038196	-0.031046	0.035141	0.034378	
fulfil	0.006534	0.028483	-0.020342	-0.000624	0.025773	0.028632	
accumulation	0.000928	-0.007928	-0.013274	-0.008226	0.036075	0.034024	
solidly	-0.009751	0.017625	0.005126	-0.008710	0.026560	0.027365	
natural	0.012698	0.028183	-0.008566	-0.016718	0.022607	0.010403	
redeem	-0.020144	-0.024476	0.003181	-0.055986	-0.010962	-0.024218	

	98	99
engaging	0.014338	0.000190
coldly	0.035130	0.010713
excellencies	0.010206	0.029503
sweetest	0.026801	0.011844
cassino	0.007945	0.032199
fulfil	0.020781	0.018778
accumulation	0.032888	0.002504
solidly	0.010209	0.025139
natural	0.004158	-0.005703
redeem	-0.037390	-0.011966


8. Measuring Similarity Between Word Pairs

9. Exploring and Visualizing your Word Embeddings using t-SNE

Sanity Check

Term: think
 Most similar words: ['deliver', 'devise', 'musical', 'murmurs', 'repetition']
 Term: thought
 Most similar words: ['rash', 'affirmative', 'plea', 'resembling', 'conceited']
 Term: mr
 Most similar words: ['overpowering', 'personal', 'urging', 'approver', 'admiring']
 Term: friend
 Most similar words: ['performers', 'enchanted', 'smiling', 'climbing', 'refrain']
 Term: love
 Most similar words: ['slept', 'repetition', 'fellows', 'procure', 'incidental']
 Term: disdain
 Most similar words: ['surpassed', 'wilfully', 'truest', 'extent', 'hating']

plt.annotate



Part B:

1. Getting the Dataset

Sanity Check:

Each instance in the training data is a list of word indices representing the words in a movie review.

Each label is 1 if that review is positive, else 0.

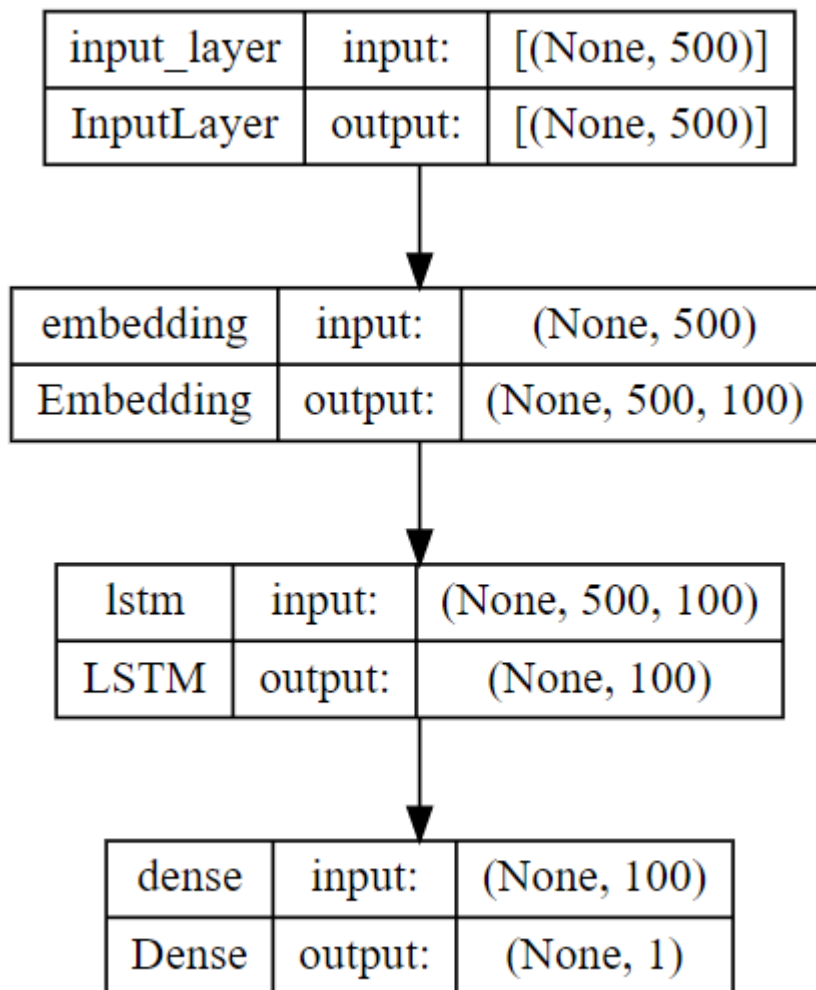
```
[ ] print('Sample review:', train_data[0])
```

Sample review: [1, 13, 21, 15, 42, 529, 972, 1621, 1384, 64, 457, 4467, 65, 3940, 3,


```
[ ] print('\n Sample label:', train_labels[0])
```

Sample label: 1

Sanity Check



4. Training the Model



Based on the accuracy plot, what do you think the optimal stopping point for your model should have been?

Where the accuracy on the test set was the highest, which is at the final epoch.

5. Evaluating the Model on the Test Data

782/782 [=====] - 27s 30ms/step - loss: 0.3362 - accuracy: 0.8698
 test_loss: 0.3361652195453644 test_accuracy: 0.8697999715805054

6. Extracting the Word Embeddings

Shape of word_embeddings: (10000, 100)

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 500, 100)	1000000
lstm (LSTM)	(None, 100)	80400
dense (Dense)	(None, 1)	101

=====
 Total params: 1,080,501
 Trainable params: 1,080,501
 Non-trainable params: 0


7. Visualizing the Reviews

<START> this film was just brilliant casting location scenery story direction everyone's really suited the part they played and you could just imagine being there robert <UNK> is an

8. Visualizing the Word Embeddings

	0	1	2	3	4	5	6	7	8	9	...	90	91	92	93	94	95	96
woods	0.034749	-0.022929	-0.042274	0.003660	0.012328	0.019674	0.047157	0.030400	0.039966	0.005881	...	-0.025695	-0.020659	-0.029948	-0.021821	-0.003663	-0.001473	-0.038691
hanging	0.023770	0.005806	-0.026436	0.023907	-0.036522	-0.003042	0.006303	0.054684	0.021125	-0.000722	...	-0.001673	0.021840	0.005051	-0.039812	-0.014943	-0.031275	-0.044489
woody	0.012521	-0.047144	-0.030312	0.001971	-0.032417	0.008225	-0.027092	-0.019926	-0.009127	0.031912	...	0.002320	-0.051157	-0.002887	-0.041049	-0.020325	0.026693	0.051277
arranged	0.056459	0.052256	0.010988	0.041696	0.042100	-0.034009	0.054902	-0.020743	-0.012765	-0.016781	...	-0.042161	0.016000	-0.027051	0.008761	0.022420	-0.030188	-0.010200
bringing	0.037957	0.017603	-0.037762	0.003273	-0.046980	-0.029365	-0.018095	0.026762	-0.019261	0.013654	...	0.036962	-0.037561	-0.001162	0.003658	0.032474	-0.008875	-0.039150
wooden	0.034470	-0.021360	0.003335	0.041142	-0.032598	0.026591	0.040554	-0.046954	-0.032404	0.025215	...	0.000053	0.031198	0.041692	0.027852	0.034170	0.008012	-0.023703
errors	0.043822	0.024764	0.047095	-0.036585	-0.034633	-0.013739	0.037724	-0.042581	-0.006961	-0.019361	...	0.043611	0.040472	-0.036316	0.026268	-0.022406	-0.001418	-0.038634
dialogs	0.008102	-0.046629	-0.037880	-0.024437	0.014995	0.050349	0.022416	-0.043939	-0.011451	0.002020	...	-0.017627	0.009910	0.038459	0.029377	0.031789	-0.049039	-0.044007
kids	-0.012254	-0.008369	-0.035314	0.030103	-0.018072	-0.021425	0.023406	0.050894	0.006287	-0.021122	...	-0.008116	-0.005464	-0.031098	-0.039549	0.030379	-0.015896	0.032091
uplifting	0.030923	0.008830	0.025629	0.024759	-0.010428	0.028655	0.034355	0.014114	0.003080	0.043952	...	0.005992	-0.012695	0.010913	-0.028248	-0.024437	0.012366	-0.035219

9. Plot your Word Embeddings using t-SNE




10. Questions

Question 1: What do you observe?
`model.summary()` is:

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 500, 100)	1000000
dropout (Dropout)	(None, 500, 100)	0
lstm_1 (LSTM)	(None, 100)	80400
dropout_1 (Dropout)	(None, 100)	0
dense_1 (Dense)	(None, 1)	101
Total params: 1,080,501		
Trainable params: 1,080,501		
Non-trainable params: 0		

Loss and Accuracy Plots;



Test Loss and Accuracy:

```
782/782 [=====] - 21s 23ms/step - loss: 0.3483 - accuracy: 0.8703
test_loss: 0.3482823073863983 test_accuracy: 0.8703200221061707
```


This model has a marginally higher accuracy than the non-dropout model but also a higher loss.

The loss on the Validation set is lowest at Epoch 4, which has a validation accuracy of 89.85%.

Question 2:

What do you Observe?


Batch Size 1:



```
782/782 [=====] - 17s 20ms/step - loss: 0.4525 - accuracy: 0.8732
test_loss: 0.45245346426963806 test_accuracy: 0.873199999332428
```

Overfits very quickly and as such the Test loss nearly doubles from epoch 2 to 4. But the Test accuracy still the highest out of the rest. The validation accuracy is highest at epoch 2 at 90.90%. Training this model was very slow.


Batch Size len(train_data):



782/782 [=====] - 23s 27ms/step - loss: 0.6880 - accuracy: 0.6456
 test_loss: 0.6879885196685791 test_accuracy: 0.6456400156021118

The Loss decreases very slowly and the test accuracy stays very low. This is most likely due to very little stochasticity in the training due to the batch size being really high. This model performs very badly.

Batch Size 32:



782/782 [=====] - 22s 27ms/step - loss: 0.4342 - accuracy: 0.8619
 test_loss: 0.434171199798584 test_accuracy: 0.861880044059753

The model overfits and the test accuracy increases after epoch 2. The Test accuracy is still very high. So, when the batch size is very high, the model doesn't have much stochasticity and as such cannot jump out of local optima. When decreasing the batch size, the amount of stochasticity increases but the training time also increases.



Part C:

1. Model 1

Model: "sequential"

Layer (type)	Output Shape	Param #
lambda (Lambda)	(None, 256, 10000)	0
global_average_pooling1d_masked (GlobalAveragePooling1DMasked)	(None, 10000)	0
dense (Dense)	(None, 16)	160016
dense_1 (Dense)	(None, 1)	17

=====
 Total params: 160,033
 Trainable params: 160,033
 Non-trainable params: 0
 =====



782/782 [=====] - 28s 34ms/step - loss: 0.4613 - accuracy: 0.7789


Note that this accuracy/loss is higher/lower than expected due to the fact I used TPU for my computations.

2. Model 2

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 256)	2560000
global_average_pooling1d_masked_1 (GlobalAveragePooling1D)	(None, 256)	0
dense_2 (Dense)	(None, 1)	257

=====
Total params: 2,560,257
Trainable params: 2,560,257
Non-trainable params: 0
=====



782/782 [=====] - 21s 25ms/step - loss: 0.3105 - accuracy: 0.8746


By using a word embedding layer, the accuracy of the model went up by 10%

3. Model 3

a. Model 3-1

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_2 (GlobalAveragePooling1DMasked)	(None, 300)	0
dense_3 (Dense)	(None, 16)	4816
dense_4 (Dense)	(None, 1)	17
=====		
Total params: 120,005,133		
Trainable params: 4,833		
Non-trainable params: 120,000,300		




782/782 [=====] - 21s 26ms/step - loss: 0.5670 - accuracy: 0.7089

Model: "sequential_3"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_3 (GlobalAveragePooling1D)	(None, 300)	0
dense_5 (Dense)	(None, 16)	4816
dense_6 (Dense)	(None, 1)	17

=====
 Total params: 120,005,133
 Trainable params: 120,005,133
 Non-trainable params: 0



782/782 [=====] - 26s 27ms/step - loss: 0.4908 - accuracy: 0.8551


So, from the above we can see that using a fine-tuning a pre-trained word embedding results in a better performance than not fine-tuning it. Interestingly, fine-tuning the pretrained word embeddings overfits very quickly. The highest val accuracy is 88.56% which actually outperforms Model 2, training our own word embedding.

b. Model 3-2

Model: "sequential_4"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
lstm (LSTM)	(None, 100)	160400
dense_7 (Dense)	(None, 1)	101

=====
 Total params: 120,160,801
 Trainable params: 120,160,801
 Non-trainable params: 0
 =====




782/782 [=====] - 24s 29ms/step - loss: 1.0130 - accuracy: 0.5714

Model: "sequential_5"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
lstm_1 (LSTM)	(None, 100)	160400
dense_8 (Dense)	(None, 1)	101

=====
 Total params: 120,160,801
 Trainable params: 160,501
 Non-trainable params: 120,000,300
 =====




782/782 [=====] - 24s 29ms/step - loss: 0.7186 - accuracy: 0.5864

Interestingly using a pre-trained word embedding with and LSTM doesn't perform too well. The non-tweakable pre-trained word embedding doesn't perform well, as expected. But surprisingly the tweakable pre-trained word embedding doesn't perform too well either. This could be due to the fact that I used a TPU to train which caused the model to overfit quickly (as the model does reach a validation accuracy of 79.59%). Another explanation could be that the learning rate is too high. The standard ADAM optimizer's learning rate may be too high for this situation, especially since the embedding is pre-trained.

4. Model 4

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
)		
global_average_pooling1d_masked (GlobalAveragePooling1D)	(None, 300)	0
dense (Dense)	(None, 100)	30100
dense_1 (Dense)	(None, 16)	1616
dense_2 (Dense)	(None, 1)	17
=====		
Total params: 120,032,033		
Trainable params: 120,032,033		
Non-trainable params: 0		




782/782 [=====] - 20s 24ms/step - loss: 0.9167 - accuracy: 0.8450

Model: "sequential_1"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_1 (GlobalAveragePooling1DMasked)	(None, 300)	0
dense_3 (Dense)	(None, 100)	30100
dense_4 (Dense)	(None, 16)	1616
dense_5 (Dense)	(None, 1)	17

=====
 Total params: 120,032,033
 Trainable params: 31,733
 Non-trainable params: 120,000,300



782/782 [=====] - 20s 25ms/step - loss: 0.5345 - accuracy: 0.7306

Both models above have a pre-trained embedding with an extra dense layer. The first model's pre-trained embedding is fine-tunable and as such results in a better performance, as expected. It achieves a final validation accuracy of 84.50%. We can see that this model starts overfitting early, with its highest validation accuracy being 88.06%. The non fine-tunable model performs much better just using an LSTM layer, as earlier This has an accuracy of 73.06%.

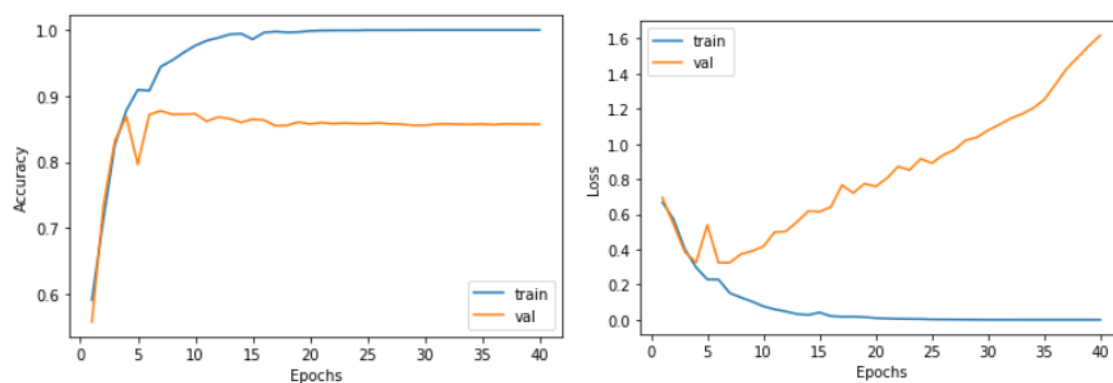
Model: "sequential_2"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_2 (GlobalAveragePooling1D)	(None, 300)	0
dense_6 (Dense)	(None, 300)	90300
dense_7 (Dense)	(None, 300)	90300
dense_8 (Dense)	(None, 100)	30100
dense_9 (Dense)	(None, 16)	1616
dense_10 (Dense)	(None, 1)	17

Total params: 120,212,633

Trainable params: 120,212,633

Non-trainable params: 0




782/782 [=====] - 20s 25ms/step - loss: 1.7369 - accuracy: 0.8431

Model: "sequential_3"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
global_average_pooling1d_masked_3 (GlobalAveragePooling1DMasked)	(None, 300)	0
dense_11 (Dense)	(None, 300)	90300
dense_12 (Dense)	(None, 300)	90300
dense_13 (Dense)	(None, 100)	30100
dense_14 (Dense)	(None, 16)	1616
dense_15 (Dense)	(None, 1)	17

=====
 Total params: 120,212,633
 Trainable params: 212,333
 Non-trainable params: 120,000,300



782/782 [=====] - 20s 25ms/step - loss: 0.5394 - accuracy: 0.7349

Both models above use the pre-trained word embedding with several hidden layers. As expected again, the fine-tunable word embedding performs better than the non-finetunable one. The fine-tunable model has a final validation accuracy of 84.31%, which is lower than the same fine-tunable model with less hidden layers. I think this is because the model overfits very quickly. Also, the non-fine-tunable model has a final validation accuracy of 73.49%, which is only marginally higher than the non-fine-tunable model with less hidden layers.

5. Model 5


a. Model 5-1

Here I use pre-trained word embeddings.

Model: "sequential"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
conv1d (Conv1D)	(None, 251, 100)	180100
global_max_pooling1d (GlobalMaxPooling1D)	(None, 100)	0
dense (Dense)	(None, 1)	101

=====
 Total params: 120,180,501
 Trainable params: 120,180,501
 Non-trainable params: 0
 =====



782/782 [=====] - 18s 22ms/step - loss: 0.3548 - accuracy: 0.8638


Here I have only used a fine-tunable Model. This has a final validation accuracy of 86.38%, which is also the highest.

b. Model 5-2

Model: "sequential_1"

Layer (type)	Output Shape	Param #
GloVe_Embeddings (Embedding)	(None, 256, 300)	120000300
conv1d_1 (Conv1D)	(None, 251, 100)	180100
conv1d_2 (Conv1D)	(None, 246, 100)	60100
global_max_pooling1d_1 (GlobalMaxPooling1D)	(None, 100)	0
dense_1 (Dense)	(None, 1)	101

=====
 Total params: 120,240,601
 Trainable params: 120,240,601
 Non-trainable params: 0



782/782 [=====] - 18s 22ms/step - loss: 0.4557 - accuracy: 0.8504

This final model has 2 Convolution layers with a fine-tunable pre-trained embedding layer. Interestingly, it doesn't perform as well (final validation accuracy of 85.04%) as only using 1 convolution layer. This is likely due to overfitting and due to the number of parameters, getting stuck in a local optimal. This could also be due to the fact that with a convolution, we are capturing the relationship in subsequences rather than the whole sequence.

Part D:**1. Implementing the Encoder**

- a. Firstly I created the Embedding lookups which just involves creating 2 embedding layers, one for the source, another for the target.
 - i. The source embedding will be configured for the source text (I.e., we use `self.vocab_source_size` for the input dimensions into the source embedding)
 - ii. I set `mask_zero=True` to ensure that padding is ignored
- b. Now that the embeddings are created, I pass the inputs of the encoder (`source_words`, and `target_words`) to the correct embedding.
 - i. At the same time, I have created and used a Dropout layer, which receives the output of the embedding layer as input
- c. I now create the encoder LSTM, which has `return_state=True` so the encoder has access to the hidden state and cell state (which are used to initialize the initial hidden states of the Training Decoder).

```

190 Task 1 encoder
191
192 Start
193 """
194 # The train encoder
195 # (a.) Create two randomly initialized embedding lookups, one for the source, another for the target.
196 print('Task 1(a): Creating the embedding lookups...')
197 embeddings_source = Embedding(self.vocab_source_size, self.embedding_size, mask_zero=True)
198 embeddings_target = Embedding(self.vocab_target_size, self.embedding_size, mask_zero=True)
199
200 # (b.) Look up the embeddings for source words and for target words. Apply dropout each encoded input
201 print('\nTask 1(b): Looking up source and target words...')
202 source_words_embeddings = Dropout(rate = self.embedding_dropout_rate)(embeddings_source(source_words))
203
204 target_words_embeddings = Dropout(rate = self.embedding_dropout_rate)(embeddings_target(target_words))
205
206 # (c.) An encoder LSTM() with return sequences set to True
207 print('\nTask 1(c): Creating an encoder')
208 encoder_lstm = LSTM(self.hidden_size, recurrent_dropout=self.hidden_dropout_rate, return_sequences=True, return_state=True)
209
210 encoder_outputs, encoder_state_h, encoder_state_c = encoder_lstm(source_words_embeddings)
211
212 End Task 1
213 """

```

2. Implementing the Decoder and the Inference Loop

- a. In this step, I get the initial states for the Inference Decoder from the final hidden states of the Training Decoder. This is passed to the `decoder_lstm`
- b. Add attention if needed by using the `decoder_attention` trained in the Training Decoder.
- c. Then, simply pass the decoder outputs with or without attention to the decoder dense layer (From Training Decoder). This is the output of the model.

```

251 | """
252 | Task 2 decoder for inference
253 |
254 | Start
255 | """
256 | # Task 1 (a.) Get the decoded outputs
257 | print('\nPutting together the decoder states')
258 | # get the initial states for the decoder, decoder_states
259 | # decoder states are the hidden and cell states from the training stage
260 | decoder_states = [decoder_state_input_h, decoder_state_input_c]
261 | # use decoder states as input to the decoder lstm to get the decoder outputs, h, and c for test time inference
262 | decoder_outputs_test, decoder_state_output_h, decoder_state_output_c = decoder_lstm(target_words_embeddings, initial_state = decoder_states)
263 | # Task 1 (b.) Add attention if attention
264 | if self.use_attention:
265 |     decoder_outputs_test = decoder_attention([encoder_outputs_input, decoder_outputs_test])
266 |
267 | # Task 1 (c.) pass the decoder_outputs_test (with or without attention) to the decoder dense layer
268 | decoder_outputs_test = decoder_dense(decoder_outputs_test)
269 |
270 | """
271 | End Task 2
272 | """

```

It looks like you forgot to detokenize your test data, which may hurt your score.

If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.

Model BLEU score: 5.53

Time used for evaluate on dev set: 0 m 10 s

Starting training epoch 8/10

240/240 [=====] - 30s 124ms/step - loss: 1.4371 - accuracy: 0.3923

Time used for epoch 8: 0 m 29 s

Evaluating on dev set after epoch 8/10:

That's 100 lines that end in a tokenized period ('.')

It looks like you forgot to detokenize your test data, which may hurt your score.

If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.

Model BLEU score: 5.66

Time used for evaluate on dev set: 0 m 9 s

Starting training epoch 9/10

240/240 [=====] - 30s 125ms/step - loss: 1.4123 - accuracy: 0.3964

Time used for epoch 9: 0 m 30 s

Evaluating on dev set after epoch 9/10:

That's 100 lines that end in a tokenized period ('.')

It looks like you forgot to detokenize your test data, which may hurt your score.

If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.

Model BLEU score: 5.72

Time used for evaluate on dev set: 0 m 10 s

Starting training epoch 10/10

240/240 [=====] - 30s 123ms/step - loss: 1.3927 - accuracy: 0.3992

Time used for epoch 10: 0 m 29 s

Evaluating on dev set after epoch 10/10:

That's 100 lines that end in a tokenized period ('.')

It looks like you forgot to detokenize your test data, which may hurt your score.

If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.

Model BLEU score: 5.67

Time used for evaluate on dev set: 0 m 10 s

Training finished!

Time used for training: 6 m 46 s

Evaluating on test set:

That's 100 lines that end in a tokenized period ('.')

It looks like you forgot to detokenize your test data, which may hurt your score.

If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.

Model BLEU score: 6.14

Time used for evaluate on test set: 0 m 10 s

Model BLEU score: 6.14

Time used for evaluate on test set: 0 m 10 s

The BLEU score is low as expected, with no attention, this model doesn't perform too well. Without Attention this model performs badly due to the fact it cannot determine which parts of the input are relevant when making a prediction. This means that each parts of the input is weighed the same even if they are not relevant.

3. Adding Attention

Assume 0-indexing on dim.

- I calculate the Luong score by doing `batch_dot` on the `encoder_outputs` and `decoder_outputs` on the second dim
- Then, softmax on the first dim.

- c. Thereafter, expand the last dim of the `luong_score`
- d. Subsequently, expand the second to last dim of the `encoder_Outputs`
- e. Then, do elementwise multiplication for the `luong_score` and `encoder_outputs`
- f. Finally, sum along the second dim (index 1).

This gives the correct `luong_score`.

```
"""
```

Task 3 attention

Start

```
"""
```

```
luong_score = K.batch_dot(encoder_outputs, decoder_outputs, (2))
luong_score = K.softmax(luong_score, 1)
luong_score = K.expand_dims(luong_score, -1)
encoder_outputs = K.expand_dims(encoder_outputs, 2)
encoder_vector = luong_score * encoder_outputs
encoder_vector = K.sum(encoder_vector, 1)
```

```
"""
```

End Task 3

```
"""
```

```
That's 100 lines that end in a tokenized period ('.')
It looks like you forgot to detokenize your test data, which may hurt your score.
If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.
Model BLEU score: 14.91
Time used for evaluate on dev set: 0 m 10 s
Starting training epoch 8/10
240/240 [=====] - 30s 124ms/step - loss: 0.9598 - accuracy: 0.5366
Time used for epoch 8: 0 m 29 s
Evaluating on dev set after epoch 8/10:
That's 100 lines that end in a tokenized period ('.')
It looks like you forgot to detokenize your test data, which may hurt your score.
If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.
Model BLEU score: 15.08
Time used for evaluate on dev set: 0 m 10 s
Starting training epoch 9/10
240/240 [=====] - 31s 128ms/step - loss: 0.9308 - accuracy: 0.5460
Time used for epoch 9: 0 m 30 s
Evaluating on dev set after epoch 9/10:
That's 100 lines that end in a tokenized period ('.')
It looks like you forgot to detokenize your test data, which may hurt your score.
If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.
Model BLEU score: 15.11
Time used for evaluate on dev set: 0 m 10 s
Starting training epoch 10/10
240/240 [=====] - 30s 127ms/step - loss: 0.9070 - accuracy: 0.5531
Time used for epoch 10: 0 m 30 s
Evaluating on dev set after epoch 10/10:
That's 100 lines that end in a tokenized period ('.')
It looks like you forgot to detokenize your test data, which may hurt your score.
If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.
Model BLEU score: 14.65
Time used for evaluate on dev set: 0 m 10 s
Training finished!
Time used for training: 6 m 58 s
Evaluating on test set:
That's 100 lines that end in a tokenized period ('.')
It looks like you forgot to detokenize your test data, which may hurt your score.
If you insist your data is detokenized, or don't care, you can suppress this message with the `force` parameter.
Model BLEU score: 14.98
Time used for evaluate on test set: 0 m 10 s
```

[101	2044	1037	3232	1997	8974	1010	1996	18726	1011	1011	1045
	2066	1996	27940	1013	24792	2621	4897	1998	1996	13675	11514	6508
	26852	1011	1011	2175	2091	2307	1012	102	8974	102	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0]				
[1	1	1	1	1	1	1	1	1	1	1	1	1
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0]				
[101	2044	1037	3232	1997	8974	1010	1996	18726	1011	1011	1045
	2066	1996	27940	1013	24792	2621	4897	1998	1996	13675	11514	6508
	26852	1011	1011	2175	2091	2307	1012	102	8974	102	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0]				
[1	1	1	1	1	1	1	1	1	1	1	1	1
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0]				

2. Basic classifiers using BERT

Layer (type)	Output Shape	Param #	Connected to
input_token (InputLayer)	[(None, 128)]	0	[]
masked_token (InputLayer)	[(None, 128)]	0	[]
tf_distil_bert_for_sequence_classification (TFDistilBertForSequenceClassification)	TFSequenceClassifierOutput(loss=None, logits=(None, 3), hidden_states=None, attentions=None)	66955779	['input_token[0][0]', 'masked_token[0][0]']
Total params: 66,955,779			
Trainable params: 66,955,779			
Non-trainable params: 0			
42/42 [=====] - 7s 91ms/step - loss: 1.0286 - accuracy: 0.8174 [1.0286388397216797, 0.817365288734436]			

Model: "Model2_BERT"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 128)]	0	[]
input_2 (InputLayer)	[(None, 128)]	0	[]
tf_distil_bert_model (TFDistilBertModel)	TFBaseModelOutput(last_hidden_state=(None, 128, 768), hidden_states=None, attentions=None)	66362880	['input_1[0][0]', 'input_2[0][0]']
global_average_pooling1d_masked (GlobalAveragePooling1DMasked)	(None, 768)	0	['tf_distil_bert_model[0][0]']
dense (Dense)	(None, 16)	12304	['global_average_pooling1d_masked[0][0]']
dense_1 (Dense)	(None, 3)	51	['dense[0][0]']
=====			
Total params: 66,375,235			
Trainable params: 66,375,235			
Non-trainable params: 0			
42/42 [=====] - 8s 93ms/step - loss: 0.6841 - accuracy: 0.8204 [0.6841117143630981, 0.8203592896461487]			

The second model with distilBERT, a dense layer and a masked pooling layer performs better than just using distilBERT for classification, but only just. Interestingly, the second model starts overfitting very quickly with the validation accuracy plateauing quickly.

For lab 4, we didn't do any sentiment analysis???

3. Advanced classifier using BERT

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 128)]	0	[]
input_4 (InputLayer)	[(None, 128)]	0	[]
tf_distil_bert_model_1 (TFDistilBertModel)	TFBaseModelOutput(last_hidden_state=(None, 128, 768), hidden_states=None, attentions=None)	66362880	['input_3[0][0]', 'input_4[0][0]']
lstm (LSTM)	(None, 100)	347600	['tf_distil_bert_model_1[0][0]']
dense_2 (Dense)	(None, 3)	303	['lstm[0][0]']
=====			
Total params: 66,710,783			
Trainable params: 66,710,783			
Non-trainable params: 0			

```
42/42 [=====] - 9s 110ms/step - loss: 0.4810 - accuracy: 0.8353  
[0.4810287654399872, 0.8353293538093567]
```

Here I used distilBERT with an LSTM layer. It outperforms the models from 2, the basic classifier. This could be because LSTMs perform a lot better at language modelling than just using a Dense Layer. But because of the use of the distilBERT model, the LSTM version, here, doesn't significantly outperform the non-LSTM versions, with only about 1-1.5% improvement in performance.